

Image Denoising with the UNet and Autoencoders

Alan Casey and Alex Racapé

November 2023

1 Introduction

In this study, we tested the U-Net and Autoencoder architectures to solve the problem of image denoising. This goal of this problem is to remove noise from an image to yield a cleaner result. U-Nets and Autoencoders make for an interesting comparison since they share some similarities in their architecture. U-Nets are commonly used for image segmentation and it centers around a sort of U shape. Convolutional and max-pooling layers reduce the dimensions of each channel while building up a greater depth of channels. Then at the bottom of the "U", transpose convolutional layers reduce the amount of channels while increasing the dimensions of each channel. The idea is to reduce an image to some core features then construct a new image from these features. Autoencoders are similar. Convolutional layers can be used for dimensionality reduction until a vector contains some core features in a lower dimensional latent space. A decoder can then reconstruct an image from the features of this latent vector. These are essentially the same process except the U-Net adds some skip connections to the network's architecture. This study contrasts the results from these two networks, and isolates the impact of these added skip connections. We constructed our own simplified networks and used the MNIST dataset to evaluate the two models' ability to denoise images both qualitatively and quantitatively.

2 Methodology

Our studies aimed to investigate the relative performance of the U-Net and Autoencoder models on the MNIST dataset. For our denoising task, we created a custom dataset object that could be used to dynamically set the amount of noise used while training and testing a network (Line 37). This dataset has an attribute *mu* which corresponds to the amount of noise in the image. We used the following equation to add noise to our image where *I* is the original image and *N* is tensor of random noise.

$$Y = (1 - \mu) * I + \mu * N$$

μ can then be interpreted as the percentage of the image that is random noise. For our experiment we defined three levels of noise. Mild noise corresponds with a μ of .1, moderate noise corresponds with a μ of .25, and severe noise corresponds with a μ of .4. We set these thresholds qualitatively based on some test plots. Figure 1 shows what these levels of noise look like with our data.




Figure 1: MNIST Data with Varying levels of Noise

Our two models (UNet and Autoencoder) were simplified versions of the U-Net architecture with and without the skip connections. They can be found on lines 69 and 145, and the model summaries show that they are essentially the same. In our implementation, the Autoencoder has slightly more weights to learn in the transpose convolutional layers since we wanted the number of channels to match up with U-Net. U-Net concatenates previous channels, so the Autoencoder needs to learn more filters in order to match this depth. However, we made sure the total number of parameters for both were as close as possible for our specified architecture: 88,641 for U-Net, and 104,521 for Autoencoder.

Since the input images from MNIST are only 28x28, we did not want to fully replicate the original U-Net. Since we were simplifying the problem by using smaller images with less channels, a reduction in channel dimension as well as a reduction in total layers makes sense. Where the classic U-Net has four main tiers with skip connections, ours has two skip connection layers. The image gets reduced from a size of 28x28 to 14x14 to 7x7, and the number of channels gets increased from 1 to 20 to 40 to 80. Then the size scales back up while the channel size scales back down resulting in a single new output image.

To test our models, we trained three versions of each network that were trained on varying levels of noise. The noisy image was the input to the network, and the model trained to output a clean image that matched the original. We trained our models for ten epochs. Then we ran several tests on the trained models. First, we plotted 64 images with the model’s predictions. Then we calculated the average MSE across 1000 test images. While training and testing, we also used timers to track the training and inference times. All training and testing was done using a T4 GPU.




Figure 2: U-Net and Autoencoder Results with Mild Noise

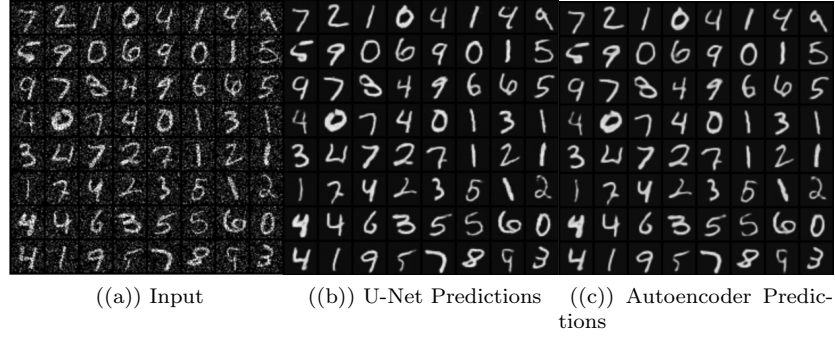


Figure 3: U-Net and Autoencoder Results with Moderate Noise

3 Results and Discussion

3.1 Qualitative Observations

For our qualitative analysis, we used 64 data points and plotted results from our models using varying levels of noise. In general, both models performed really well with results that closely resembled the clean original images. While the results were very similar, the results from the Autoencoder seemed slightly more blurry, especially for the lower noise levels. Both models struggled more with the severe level of noise, yet the results look very similar. The only noticeable difference is the six in the bottom right hand corner, where Autoencoder’s result looks more like a pretzel. With this difference U-Net may have done slightly better, but it is hard to say conclusively through a purely qualitative analysis.

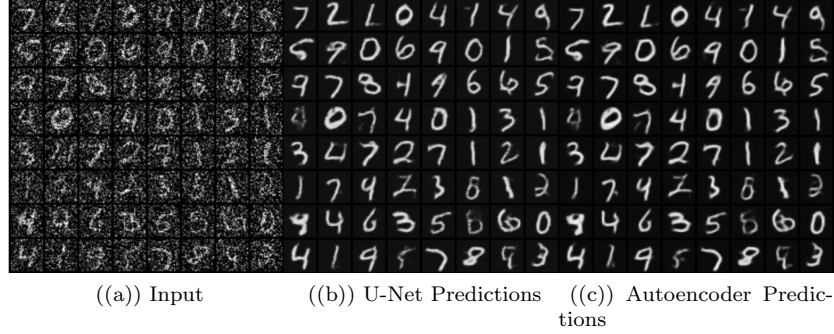


Figure 4: U-Net and Autoencoder Results with Severe Noise

Table 1: Model Performance Comparison

Model	Noise	Average MSE	Training Time	Inference Time
U-Net	Mild	0.00095	1m 4s	1.3ms
	Moderate	0.0047	1m 8s	1.3ms
	Severe	0.014	1m 4s	1.4ms
Autoencoder	Mild	0.0020	1m 7s	2.6ms
	Moderate	0.0055	1m 11s	1.4ms
	Severe	0.015	1m 6s	2.0ms

3.2 Quantitative Observations

In order to get more accurate measures that could help us in our analysis, we decided to measure differences from a quantitative perspective. We decided to measure both our UNet and Autoencoder networks on their resulting MSEs (Mean Squared Error Loss) after ten epochs, their total training time, and their inference time.

The Mean Squared Error Loss Function creates a criterion that measures the mean squared error between each pixel from the original image to the output image from the neural network. This measures how far off the de-noised image was from the original image.

As shown in Table 1, our UNet model’s average MSE outperformed the Autoencoder’s average MSE in all three noise categories, with the most significant out-performance being in the Mild-noise category. One of the reasons for this could be that, even though the Autoencoder has more weights it can learn, the UNet employs skip-connections. This helps preserve information, and these skip connections allow for a direct flow of information from the early layers (which extract low-level features) to the later layers (which contain higher-level features) of the UNet network. These added channels contribute some information that is lost for the Autoencoder.

The total training time was also faster for UNet—this is most likely the

result of having fewer parameters. Skip connections often result in fewer parameters compared to fully connected or densely connected networks. This is because skip connections allow the re-utilization of features from lower layers instead of creating entirely new sets of parameters. This is another potential reason as to why the UNet outperformed the Autoencoder in its average MSE: fewer parameters can help prevent the overfitting of data while reducing the computational cost (which also reduces the total inference time for each image).

4 Conclusion

Our study aimed to compare the performance of both the UNet and Autoencoder architectures for the task of image denoising. We did this using the MNIST dataset of handwritten numbers. Both models showed promising results in reducing noise and restoring images to their cleaner versions. However, there were notable differences between the two models, with our UNet architecture demonstrating a clear advantage.

Qualitatively, when analyzing the output images generated by both models, we observed that both U-Net and Autoencoder produced impressive denoised images which closely resembled the original images. However, the Autoencoder’s results were slightly more blurry, especially for images with lower levels of noise. However, it was still challenging to make a definitive conclusion based solely on a qualitative analysis.

To gain a more quantitative perspective, we measured the Mean Squared Error (MSE) loss for both models after training on images with varying levels of noise. The results revealed that UNet consistently outperformed the Autoencoder in terms of average MSE, with the most significant difference observed in the mild noise category. This superior performance could be attributed to U-Net’s use of skip connections, which facilitate the preservation of fine-grained details and transfer information between different network layers. This allows U-Net to handle the denoising task more effectively.

Furthermore, our UNet exhibited shorter training times and faster inference times compared to the Autoencoder, likely due to its fewer parameters. Skip connections in our UNet can lead to more efficient training and reduced computational cost, which can be advantageous in practical applications.

In summary, our study highlights the effectiveness of both the U-Net and Autoencoder architectures for image denoising. However, U-Net, with its skip connections and fewer parameters, outperformed the Autoencoder in terms of denoising accuracy and computational efficiency. These findings provide valuable insights for researchers and practitioners working on image denoising tasks and related applications.

5 Appendix

```
1  # -*- coding: utf-8 -*-
2  """lab6
3
4  Automatically generated by Colaboratory.
5
6  Original file is located at
7      https://colab.research.google.com/drive/1V8xuNuqp8dkIQoRJZR2Hc_2fWvWMm2rd
8  """
9
10 # Project 6: U-Net and Autoencoders
11
12 # Imports
13 import torch
14 from torchvision.datasets import MNIST
15 from torch.utils.data import Dataset, DataLoader
16 from torchvision import transforms
17 import torch.nn as nn
18 from torchsummary import summary
19 from torch.optim import Adam
20 import numpy as np
21 import matplotlib.pyplot as plt
22 import time
23
24 device = "cuda" if torch.cuda.is_available() else "cpu"
25 print(device)
26
27 # Global noise values
28 MILD = .1
29 MEDIUM = .25
30 SEVERE = .4
31
32 # Load MNIST data
33 mnist_train = MNIST('~/.data', train=True, download=True)
34 mnist_test = MNIST('~/.data', train=False, download=True)
35
36
37 class MNISTDataset(Dataset):
38     """Noisy wrapper for MNIST Dataset
39
40     We can dynamically set the noise using the mu attribute
41     X's are noisy images, and Y's are the clean original images
42     Images are (1, 28, 28)
43     """
44
45     def __init__(self, x, y, mu):
46
47         # Reformat input image
```

```

48     x = x.view(-1, 1, 28, 28)
49     normal_image = x.float()/255
50
51     # Add the noise to the image, but keep vals in [0, 1]
52     noise = torch.randn_like(normal_image, dtype=torch.float)
53     noisy_image = (1 - mu) * normal_image + (mu) * noise
54     noisy_image = torch.clamp(noisy_image, 0, 1)
55
56     # Save as input and ground-truth
57     self.x, self.y = noisy_image, normal_image
58
59     def __getitem__(self, ix):
60         return self.x[ix].to(device), self.y[ix].to(device)
61
62     def get_group(self, size):
63         """Get subset with `size` images"""
64         return self.x[0:size].to(device), self.y[0:size].to(device)
65
66     def __len__(self):
67         return len(self.x)
68
69     class UNet(nn.Module):
70
71         """
72         Our UNet is a neural network that reduces an images to core features
73         (while reducing the size of the image and increasing the channels),
74         and then increases it back to its original size.
75
76         Since the input images from MNIST are size 28x28, we did not want to
77         fully replicate the original UNet. Because we are simplifying the problem,
78         a reduction in channel dimension makes sense here, as well as a reduction
79         in total layers.
80
81         We decided our architecture would look like this:
82         - 1 CL (Convolutional Layer) followed by a Max Pool Layer 2 times, then
83         - A bottom CL followed by a TCL (Transpose CL) 2 times, followed by
84         - Two CL (the last one reducing the dimension size back to 1)
85
86         The image gets reduced from a size of 28x28 -> 14x14 -> 7x7, the channel
87         size gets increased from 1 -> 20 -> 40 -> 80, and then the size scales
88         back up and the channel size scales back down.
89
90         One key aspect of the UNet is that before the MaxPool layers, a copy of
91         the outputs get saved so it can be concatenated later in the re-scaling side.
92         """
93
94         def __init__(self, channels=1, output_classes=10):
95             super().__init__()
96             self.loss_func = nn.MSELoss()
97

```

```

98         # Convolution Layers
99         self.ConvLayer1 = self.ConvLayer(1, 20, 3)
100         self.MaxPool = nn.MaxPool2d(2)
101         self.ConvLayer2 = self.ConvLayer(20, 40, 3)
102
103         # Bottom Layer
104         self.ConvLayer3 = self.ConvLayer(40, 80, 3)
105
106         # Transpose Convolution Layers
107         self.TransposeLayer1 = self.TransposeConvLayer(80, 40, 2)
108         self.ConvLayer4 = self.ConvLayer(80, 40, 3) # Concatenate ConvLayer2 result
109         self.TransposeLayer2 = self.TransposeConvLayer(40, 20, 2)
110         self.ConvLayer5 = self.ConvLayer(40, 20, 3) # Concatenate ConvLayer1 result
111
112         # Flatten to final image
113         self.ConvOneByOne = nn.Conv2d(20, 1, kernel_size=1)
114
115     def ConvLayer(self, in_ch, out_ch, kernel_size=3):
116         sequence = nn.Sequential(
117             nn.Conv2d(in_ch, out_ch, kernel_size, 1, 1),
118             nn.ReLU())
119         return sequence
120
121     def TransposeConvLayer(self, in_ch, out_ch, kernel_size=2):
122         sequence = nn.Sequential(
123             nn.ConvTranspose2d(in_ch, out_ch, kernel_size, 2),
124             nn.ReLU())
125         return sequence
126
127     def forward(self, x):
128         a = self.ConvLayer1(x)
129         x = a
130         b = self.ConvLayer2(self.MaxPool(x))
131         x = b
132         x = self.TransposeLayer1(self.ConvLayer3(self.MaxPool(x)))
133         x = self.TransposeLayer2(self.ConvLayer4(torch.cat([b,x], dim=1)))
134         return self.ConvOneByOne(self.ConvLayer5(torch.cat([a,x], dim=1)))
135
136     def num_parameters(self):
137         return sum(p.numel() for p in self.parameters() if p.requires_grad)
138
139
140     # Create a UNet instance and visualize parameter grid
141     unet = UNet().to(device)
142     print(f"Total parameters for UNet: {unet.num_parameters()}")
143     summary(unet, (1, 28, 28))
144
145     class Autoencoder(nn.Module):
146
147         """

```



```

148     Our Autoencoder network is essentially the same thing as our UNet.
149     However, we do not employ skip connections (we don't save the output
150     of the Convolutional Layers in order to do concatenations later).
151
152     The autoencoder employs the same architecture and size reduction/
153     scaling sizes as our previous UNet architecture.
154 """
155
156 def __init__(self, channels=1, output_classes=10):
157     super().__init__()
158     self.loss_func = nn.MSELoss()
159
160     self.layers = nn.Sequential(self.ConvLayer(1, 20, 3),
161                                 nn.MaxPool2d(2),
162                                 self.ConvLayer(20, 40, 3),
163                                 nn.MaxPool2d(2),
164                                 self.ConvLayer(40, 80, 3),
165                                 self.TransposeConvLayer(80, 40, 2),
166                                 self.TransposeConvLayer(40, 20, 2),
167                                 nn.Conv2d(20, 1, kernel_size=1))
168
169 def ConvLayer(self, in_ch, out_ch, kernel_size=3):
170     return nn.Sequential(nn.Conv2d(in_ch, out_ch, kernel_size, 1, 1),
171                           nn.ReLU())
172
173 def TransposeConvLayer(self, in_ch, out_ch, kernel_size=2):
174     return nn.Sequential(nn.ConvTranspose2d(in_ch, in_ch, kernel_size, 2),
175                           nn.ReLU(),
176                           nn.Conv2d(in_ch, out_ch, 3, 1, 1),
177                           nn.ReLU())
178
179 def forward(self, x):
180     return self.layers(x)
181
182 def num_parameters(self):
183     return sum(p.numel() for p in self.parameters() if p.requires_grad)
184
185
186 # Create an autoencoder instance and visualize parameter grid
187 autoencoder = Autoencoder().to(device)
188 print(f"Total parameters for Autoencoder: {autoencoder.num_parameters()}")
189 summary(autoencoder, (1, 28, 28))
190
191 def train_batch(model, X_train, Y_train, opt):
192     opt.zero_grad() # Flush memory
193     pred = model(X_train) # Get predictions
194     batch_loss = model.loss_func(pred, Y_train) # Compute loss
195     batch_loss.backward() # Compute gradients
196     opt.step() # Make a GD step
197     return batch_loss.detach().cpu().numpy()

```

```

198
199
200 def train(model, train_dl, epochs, optimizer):
201
202     loss_history = []
203     start = time.time()
204     for epoch in range(epochs):
205
206         print(f"Running Epoch {epoch + 1} of {epochs}")
207         epoch_losses = []
208         for i, batch in enumerate(train_dl):
209             x, y = batch
210             x, y = x.to(device), y.to(device)
211             batch_loss = train_batch(model, x, y, optimizer)
212             epoch_losses.append(batch_loss)
213
214         epoch_loss = np.mean(epoch_losses)
215         loss_history.append(epoch_loss)
216
217     end = time.time()
218     training_time = end - start
219     return loss_history, training_time
220
221
222 # Train the models
223 unet_models = []
224 unet_losses = []
225 unet_times = []
226 autoencoder_models = []
227 autoencoder_losses = []
228 autoencoder_times = []
229 noise_levels = [MILD, MEDIUM, SEVERE]
230 for noise in noise_levels:
231
232     # Create dataset and loader for training and testing
233     train_dataset = MNISTDataset(mnist_train.data, mnist_train.targets, noise)
234     train_dl = DataLoader(train_dataset, batch_size=128, shuffle=True)
235
236     # Plot a sample of the noisiness
237     normal, noisy = train_dataset[1]
238     plt.imshow(normal.cpu()[0, :, :])
239     plt.show()
240     plt.imshow(noisy.cpu()[0, :, :])
241     plt.show()
242
243     # Create Models
244     unet = UNet().to(device)
245     autoencoder = Autoencoder().to(device)
246
247     # Train competing models

```

```

248     num_epochs = 5
249     lr = .001
250     u_opt = Adam(unet.parameters(), lr=lr)
251     a_opt = Adam(autoencoder.parameters(), lr=lr)
252     u_loss_history, u_train_time = train(unet, train_dl, num_epochs, u_opt)
253     a_loss_history, a_train_time = train(autoencoder, train_dl, num_epochs, a_opt)
254
255     # Save Models
256     unet_models.append(unet)
257     unet_losses.append(u_loss_history)
258     unet_times.append(u_train_time)
259     autoencoder_models.append(autoencoder)
260     autoencoder_losses.append(a_loss_history)
261     autoencoder_times.append(a_train_time)
262
263     # Test the models
264     !pip install -q torch_snippets
265
266     from torchvision.utils import make_grid
267     from torch_snippets import show
268
269
270     def plot_loss(loss_history):
271         plt.plot(loss_history)
272         plt.title("Loss Over Time")
273         plt.xlabel("Epochs")
274         plt.ylabel("Loss Value")
275         plt.show()
276
277
278     def plot_comparison(model, x, y):
279
280         # Plot noisy input
281         true_images = x
282         true_grid = make_grid(true_images, nrow=8, normalize=True)
283         show(true_grid.cpu().detach().permute(1,2,0), sz=5)
284
285         # Plot a grid of our predictions
286         predicted_images = model(x).data.cpu().view(64, 1, 28, 28)
287         prediction_grid = make_grid(predicted_images, nrow=8, normalize=True)
288         show(prediction_grid.cpu().detach().permute(1,2,0), sz=5)
289
290         # Plot a grid of real images
291         true_images = y
292         true_grid = make_grid(true_images, nrow=8, normalize=True)
293         show(true_grid.cpu().detach().permute(1,2,0), sz=5)
294
295
296     def calculate_mse(model, x, y):
297

```

```

298     # Calculate MSE per pixel across all images pairwise
299     loss = nn.MSELoss()
300     start = time.time()
301     pred = model(x)
302     end = time.time()
303     output = loss(pred, y)
304
305     return output, end - start
306
307
308 def run_tests(model, dataset):
309
310     # Calculate quantitative stats
311     x, y = dataset.get_group(len(dataset))
312     mse, test_time = calculate_mse(model, x, y)
313     print(f"MSE: {mse}")
314     print(f>Inference time: {test_time}")
315
316     # Plot for qualitative differences
317     x, y = dataset.get_group(64)
318     plot_comparison(model, x, y)
319
320
321 # Run through our tests on the saved models
322 for i, noise in enumerate(noise_levels):
323
324     test_dataset = MNISTDataset(mnist_test.data, mnist_test.targets, noise)
325
326     print(f"-- Starting U-Net tests - Noise: {noise}")
327     print(f"Training time: {UNET_times[i]}")
328     run_tests(UNET_models[i], test_dataset)
329     print(f"-- Starting Autoencoder tests - Noise: {noise}")
330     print(f"Trainint time: {autoencoder_times[i]}")
331     run_tests(autoencoder_models[i], test_dataset)

```